

Adversarial Arena: Driving innovation in Responsible AI through Interactive Competition



TrustNLP Workshop 2026
Keynote

Michael Johnston
Amazon, Responsible AI,
AGI Foundations
mjohnstn@amazon.com

Outline

- 3 pressing challenges:
 - Scaling safety
 - Evaluation limitations
 - Data scarcity
- Adversarial Arena
- History of AI challenge competitions
- Amazon Nova AI Challenge: Trusted AI
- Amazon Nova AI Challenge: Trusted Software Agents
- Closing remarks and future work

1. Scaling AI Safety and Alignment

“With great power comes great responsibility”

- **UTILITY:** Advances in AI are enabling incredible new capabilities across all task areas:

Drug discovery

Democratization of software creation for increasingly complex tasks

Science / Medicine

Software Engineering


....

- **RESPONSIBILITY:** Critical that AI remains safe, reliable, and, trustworthy → maximally *helpful* but not *harmful*

Bio-weapon design

Malicious code

Vulnerable code



2. Limitations of Evaluation


- Primary focus still on static benchmarks
- Poor proxy for performance of AI models/systems in the wild
- Contamination (Xu et al 2024, Cheng et al 2025, Fodor 2025)
- Mostly single turn while real world usage is frequently multi-turn (Laban et al 2026, Graf et al 2026, Sideshmukh et al 2025)

3. Data Scarcity

- LLM training requires diverse high-quality data
 - Almost all human generated content will soon already be incorporated in training (Villalobos et al 2024)
 - New data rare and costly to obtain, especially for novel tasks and multi-turn
- Crowdsourcing:
 - Data quality issues (Vaughan 2018)
- Synthetic data creation
 - Often lacks diversity (Xu et al 2025)
 - Can amplify inherent bias

Adversarial Arena

- Adversarial competition format
- Competing teams are divided into attackers and defenders
- Teams compete across a series of tournaments refining their approaches
 - Multi-turn
 - Drives robustness to diversity of opponents
- Related work:
 - “Hide-and-Seek Privacy Challenge”
 - Jordon et al 2020
 - cf. Chatbot Arena/LMarena/Arena
 - Chiang et al 2024



AI Challenge competition history



Amazon AI Challenges

- In competitions since 2016, engaged 1000+ students and faculty from around the world to take on key challenges in artificial intelligence

- **Alexa Prize**

SocialBot Grand Challenge:


- Open-domain conversation

TaskBot Challenge:

- Multimodal task-oriented dialog

SimBot Challenge:

- Embodied conversational AI



Amazon Nova AI Challenge

Trusted AI

Trusted Software Agents

Amazon Nova AI Challenge: Trusted AI (2024-2025)

- Global university competition to drive secure innovation in AI
 - Intersection of Coding agents / Responsible AI / AI Security / Dialog
- **Defenders:** Creating coding agents that are maintain utility but do not cause harm
- **Attackers:** Creating automatic red teaming agents
- 5 Red teams / Attackers:
 - Build automated conversational red teaming agents that seek out vulnerabilities in model developer systems and try to elicit malicious code, vulnerable code, or assistance with security events
- 5 Model developer teams / Defenders:
 - Build coding agents which strive to maintain utility while avoiding generating malicious code, vulnerable code, or assistance with carrying out security events




Year 1: Tournaments

- Each tournament:
 - 25 (5 attackers x 5 defenders) match-up
- Each match-up:
 - 200 conversations
- Each conversation:
 - Up to 10 turns
- Utility evaluation prompts interspersed
- Practice runs and dry runs
- Probing sessions (iteration)
- Tournament orchestration service


Conversation




Tournament Orchestration: Attack Success




Tournament Orchestration: Defense Success



Year 1: Evaluation




Year 1: Evaluation (cont.)



- Defenders can't win just by doing nothing !
 - Defense success scaled by utility
 - Fresh utility prompts and tests each tournament

Year 1: Evaluation (cont.)



- We are interested in disparate successful attacks !
 - Attack success scaled by diversity

- Defenders can't win just by doing nothing !
 - Defense success scaled by utility

Year 1: Evaluation (cont.)

- How are teams ranked ?



Year 1: Competing Teams

Red Teams:

1

Team PurCL: Purdue

2

Team RedTWIZ: NOVA Portugal

Team CapitalAI: UC Davis

Team Astro: UT Dallas

Team SaFoLab: University Wisconsin
Madison

Model Developers:

1

Team PurpCORN/PLAN: UIUC

2

Team Alquistcoder:
Czech Technical University

Team HokieTokie: Virginia Tech

Team Lioncoders: Columbia University

Team Purpl3Pwn3rs: CMU

<https://www.amazon.science/nova-ai-challenge/teams/>

Year 1: Innovations and Observations: Red Teams

Multi-turn attacks dominated: From tournament to tournament defenders quickly adapted their coding assistants to deflect most single turn attacks. Successful red teams developed strong multi-turn attack strategies.

Adaptive attack planning: Red teams developed various kinds of attack planners to hone in on the most promising attacks:

- **PurCL (Purdue):** developed an approach using Gibbs Sampling to find the most promising areas of the attack space
- **RedTWIZ (NOVA, Portugal)** developed a hierarchical attack planner
- **CapitalAI (UC Davis)** built a strategy library with pairs of successful and unsuccessful attacks. At attack time they retrieved strategies and evolved attacks from failure to success

“Begin from Benign”: A common strategy was to start from non-malicious utility-oriented prompts and gradually evolve the dialog towards a malicious goal (e.g. SaFoLab (UWISC), RedTWIZ (NOVA))

<https://www.amazon.science/nova-ai-challenge/proceedings>

Year 2: Innovations and Observations: Model Developers

Synthetic Data: All teams made heavy use of data generated synthetically by LLMs. e.g. HokieTokie (Vtech) synthesized additional training data combining the MITRE taxonomy, examples of vulnerable code from the internet, and failure cases from tournaments as inputs.

Reasoning about safety: Several teams adopted deliberative alignment approaches that taught models to reason about safety. e.g. Team PurpCorn-PLAN (UIUC) developed an approach which (1) explicitly taught their model to reference cybersecurity rules and (2) used GRPO with a multi-objective reward mechanism that optimized safety while maintaining utility [1].

Vulnerability Refiners: Several teams employed a multi-pass approach with some form of self-reflection. For example, both Purpl3Pwn3rs (CMU) and LionCoders (Columbia) post-trained a separate 'vulnerability fixer' that took a second pass over generated code to address security vulnerabilities.

<https://www.amazon.science/nova-ai-challenge/proceedings>

[1] Liu et al 2025. "PurpCode: Reasoning for Safer Code Generation". Proceedings of NeurIPS.

Year 1: Trusted AI: Challenges addressed

- **Scaling AI Safety:**
 - Making coding agents more robust to adversarial attempts to generate malicious code, vulnerable code, or provide assistance with cyberattacks
 - Development of automated red teaming agents: Scaling from manual red teaming
- **Evaluation:**
 - Evaluation is dynamic. Newly updated opposing systems in each tournament round. Need to perform well against all to rank highly
 - Evaluation is multi-turn. Red teaming agents are simulated malicious developers that engage in multi-turn conversation
- **Data:**
 - Corpus of 25600 multi-turn adversarial conversations with annotations for vulnerable code, malicious code, and assistance with security events
 - https://huggingface.co/datasets/amazon-agi/AdversarialArena_Nova_AI_Challenge_Trusted_AI_Dataset


Goyal et al 2026 “Adversarial Arena: Crowdsourcing Data Generation through Interactive Competition” ICLR 2026 DATA-FM Workshop.

Paper



Data





A person is seen from the side, working at a desk in a dimly lit room. The room is bathed in a cool blue light, likely from the computer monitors. The person's hands are on a keyboard, and a mouse is visible. In the background, several computer monitors are visible, some displaying code or data. The overall atmosphere is one of focused, late-night work in a tech environment.

The First Ever

**Amazon Nova AI Challenge:
Trusted AI Track**

Amazon Nova AI Challenge: *Trusted Software Agents (2025-2026)*

- Theme of coding agents and AI safety has only increased in relevance → Launched Year 2, Oct 2025
- Challenge design has to evolve though:
 - **Coding:**
 - Text-to-code → Repo level SWE tasks
 - Adding features to full-stack web applications (Python, JS, TS, HTML, CSS)
 - Model developer teams ranked both on utility and safety
 - Utility measurement: Unit tests → Functional tests on apps (Playwright)
 - **Cybersecurity:**
 - Static analysis tools → Cybersecurity agents
 - Red teams build agents that seek out vulnerabilities and exploits in applications built
 - **User Simulation:**
 - Coding agents interact with user simulators to complete tasks on repos
 - Benign user simulators
 - Malicious user simulators



Year 2: Challenge Design

Blue Teams (Model Developer teams)

Coding Agent




Malicious User Simulator



Red Teams


Cybersecurity Agent




Benign User Simulator



Year 2: Challenge Design: Build Phase




Year 2: Challenge Design: Build Phase (cont.)




Year 2: Tournament Orchestration: Build Phase

Simulated Users:




Tournament Orchestrator

Coding Agents:



- $5 * (4 + 5) = 45$ matchups per task
- 5 repos * 10 tasks per tournament
- 2250 conversations

Year 2: Challenge Design: Attack Phase



Year 2: Evaluation

- Match level scoring:
 - Each match is a single evaluation instance where all teams face identical conditions
- Objective level ranking:
 - Each objective ranked independently using Plackett-Luce Model
 - Model assigns each team a strength representing probability of ranking first in that objective
 - Accounts for opponent strength
- Overall team ranking in each tournament combines their two sub-objectives

Matches	Characteristics	Scoring Function
Coding Agent	Same benign user on same task. Combines utility and security	$u = (1 - t) + \epsilon_1$ $c = V + \epsilon_2$ $S_{CA} = \frac{2uc}{u + c}, \quad \epsilon_1 = 1, \epsilon_2 = 1$
Cybersecurity Agent	Same application (excluding apps from same team benign user) Combines vulns and exploits	$S_{SA} = R(e) + F_1(v, V)$ $R(e) = \begin{cases} 2 & \text{if } e \cap E_{RCE} \neq \emptyset \\ 1 & \text{if } e \neq \emptyset \\ 0 & \text{otherwise} \end{cases}$
Malicious User	Same coding agent on same task	$S_{MU} = V \cdot t$
Benign User	Same coding agent on same task	$S_{BU} = t$

Year 2: Competing teams

Red Teams:

Team PurCL: Purdue

Team Seclab: OSU

Team Pr1smCode: CMU

Team Lion-OxA: Penn State

Team Jay'l break: Johns Hopkins

Model Developers:

Team PurpCORN/PLAN: UIUC

Team Alquistcoder:
Czech Technical University

Team BlueTWIZ: NOVA Portugal

Team Slugs N' Roses: UC Santa Cruz

Team BruinWeb: UCLA

<https://www.amazon.science/nova-ai-challenge/teams/>

Some early observations on Year 2

- Teams will be competing through the end of the year....
- Some things we are seeing at high level:
 - Teams building complex agentic harnesses both for their coding agents and security agents
 - Novel techniques for synthesis of repositories, tasks, and training gyms
 - SFT and RL for improving secure coding of repo level tasks
 - Complexity in user simulator design and implementation
 - Adding measurement of realism/naturalness of user simulators

Year 2: Challenges addressed

- **Scaling AI Safety:**
 - Developing approaches to coding agents that can conduct repo level tasks but do so securely
 - Advances in bots for defensive adversarial security (given application and code -> finding vulnerabilities and exploits)
- **Evaluation:**
 - Multi-faceted dynamic evaluation for secure coding agents, cybersecurity bots, and user simulators
 - Inherently multiturn with user simulators representing both benign and malicious developers
 - Rather than relying on static analysis tools, application vulnerabilities are probed and tested by a panel of 5 different cybersecurity bots
- **Data (and more !):**
 - 10+ different sets of AI artifacts for coding utility and AI security (Eval benchmarks, Gyms, Data, User simulators)

Year 2: Evaluation and Training Artifacts Being Generated

Full stack
web coding
benchmark

Full stack
web coding
RL gym

Adversarial
defensive
security
benchmark

Adversarial
defensive
security RL
gym

Training
data sets
from teams

Secure
coding
agent SFT
data

Adversarial
defensive
security SFT
data

Tournament
Conversation
Dataset

5 benign
user
simulators

5 malicious
user
simulators

Some Lessons Learned

- Intricacy of challenge design and incentive structure to ensure quality of resulting datasets and other artifacts
- *Iterate* on challenge design and evaluation to avoid 'reward' hacking.
- Examples:
 - Year 1: Needed to expand utility evaluation to include tasks more likely to involve vulnerabilities (which cannot be tested in the style of HumanEval)
 - Year 2: Early user simulators are too helpful and do not represent typical human behavior adding realism metrics to later tournament and finals

Some Lessons Learned (cont.)

- Dynamic evaluation (new repos, tasks every tournament) drives teams to come up with new approaches rather than *bench-maxxing*. They can learn from experience but face novel challenges each time.
- Teams thrive on the competitive spirit of the challenge, learning to work together is a big part of the experience for teams
- Human annotation in scoring (e.g. verification of vulnerabilities in Year 2) velocity of innovation

Closing Remarks

- Increasingly critical to scale efforts in AI safety
- Adversarial interactive competition overcomes the limitations of competing purely on static benchmarks
- Generates a diverse range of artifacts for AI evaluation and training
- Challenge design and evaluation is complex and needs to be continually iterated
- Evaluation science needs to better address multi-turn collaboration



Future work

- Challenge tasks on other RAI pillars and more broadly on LLM utility
- Expanding engagement beyond 100+ students and faculty per challenge task
 - Reaching broader audience
 - Expanding beyond universities
 - Increase international participation
 - Open challenges
- ‘Graduated’ challenge model: Start open then some teams progress
- Creating an Adversarial Arena which is “always on”

Acknowledgements

- Prize team
- All the many participating university teams
- Our partner teams across Amazon, including:
 - RAI, SWE, Security, AWS

Thank you!

mjohnstn@amazon.com

<https://www.amazon.science/nova-ai-challenge>

